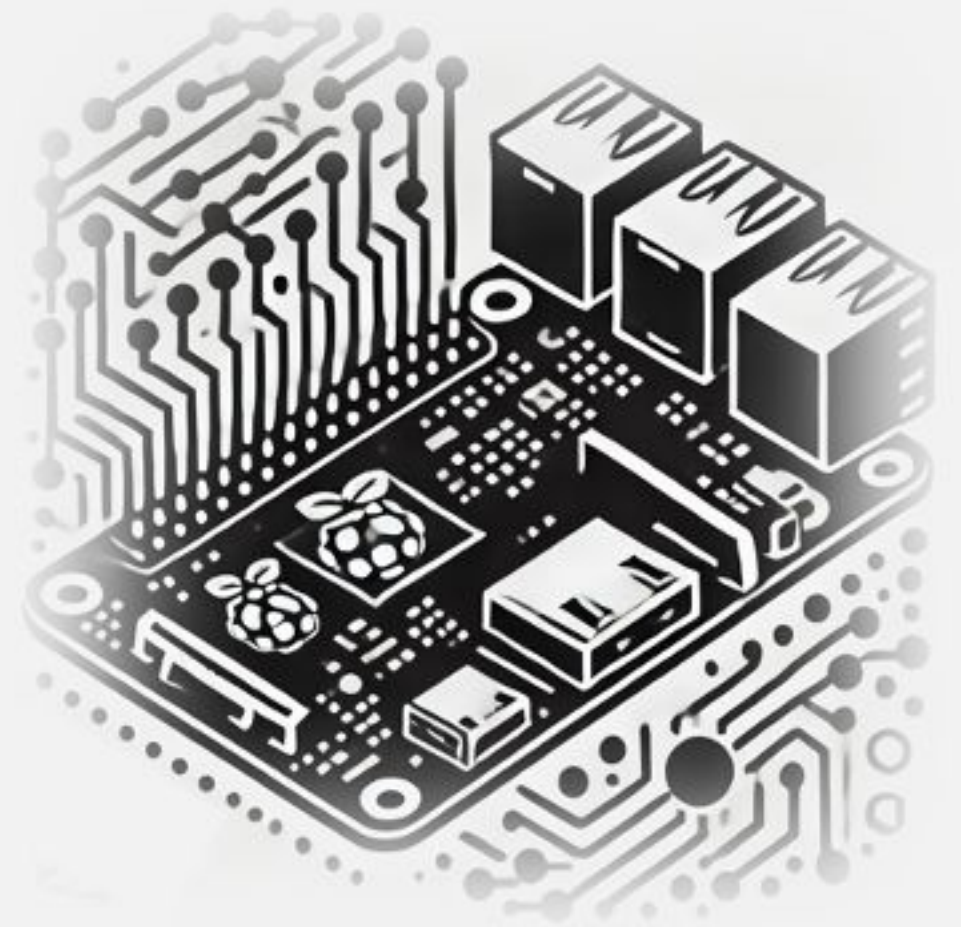


IESTI05 – Edge AI

Machine Learning System Engineering

4. Image Classification: Introduction



Computer Vision Recognition Tasks

Image Classification (Multi-Class Classification)




Cat: 70%



Dog: 80%

Object Detection Multi-Label Classification + Object Localization



Computer Vision Recognition Tasks

Instance Segmentation

Each **pixel** in an image IS CLASSIFIED into a predefined category.



Pose Estimation

Key points (or landmarks) on the object, such as joints on a human body are detected



Image Classification

What is Image Classification?

Definition

Image classification is a fundamental task in computer vision that involves **categorizing an image into one of several predefined classes**.

Key Characteristics

- Assigns a **single label to the entire image**
- Based on visual content analysis
- Uses machine learning algorithms
- Mimics human visual perception



Cat: 70%



Dog: 80%

Real-World Applications

Healthcare

Medical image analysis, X-ray and MRI diagnostics

Agriculture

Crop health monitoring, disease detection

Automotive

Autonomous vehicles, road sign recognition

Security

Surveillance systems, threat detection


Retail

Visual search, inventory management


Environmental

Satellite imagery analysis, monitoring

Image Classification Inference Pipeline




Input Image



Pre-Process



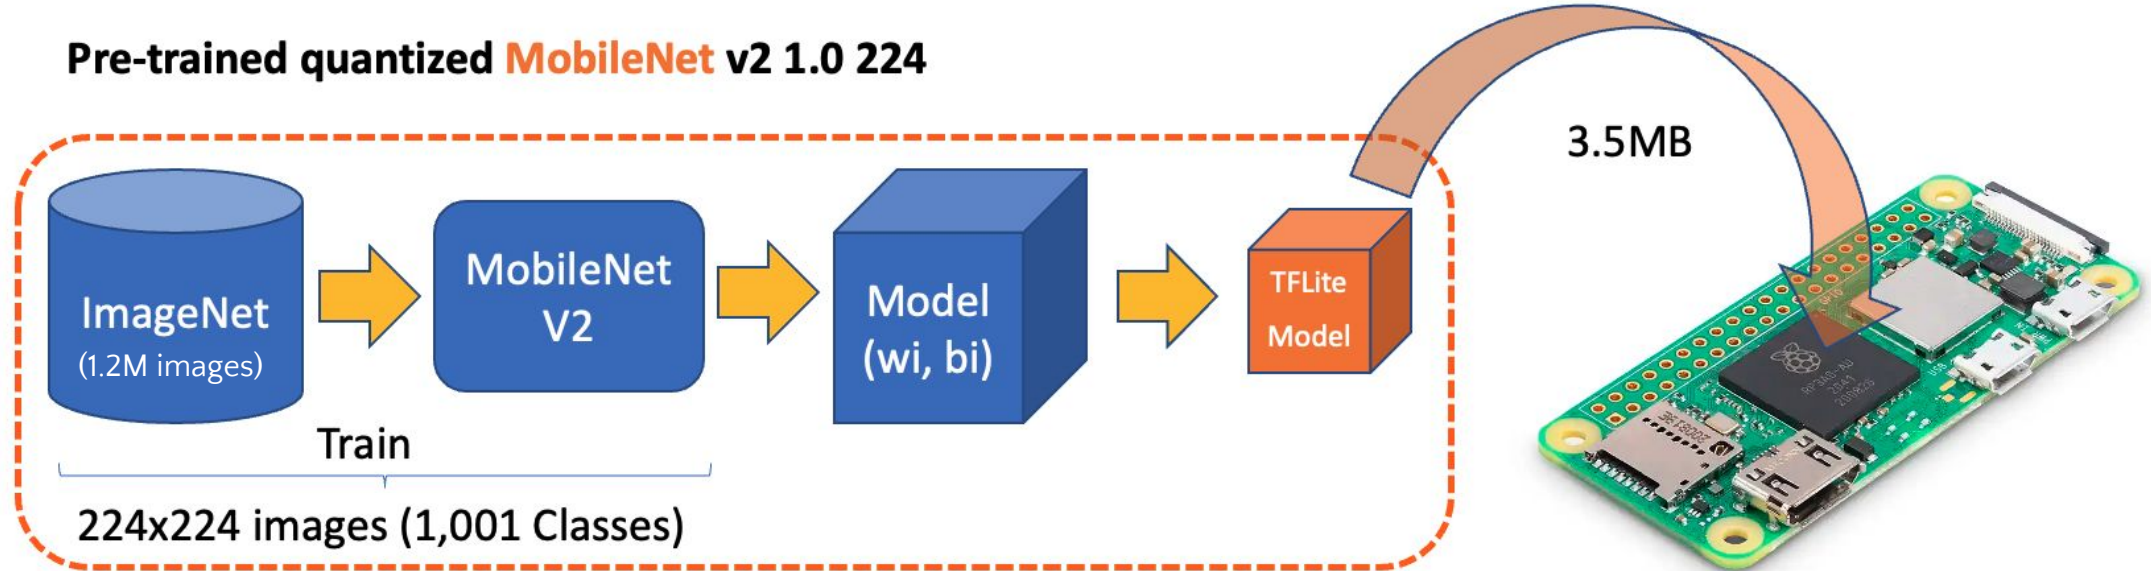
Classification
Model



Predict Label

The MobileNet V2 model

Pre-trained quantized **MobileNet v2 1.0 224**




Setting up a Virtual Environment

Activate the environment:

```
source ~/tflite_env/bin/activate
```

To **exit** the virtual environment, use:

```
deactivate
```



TensorFlow Lite Setup

What is TensorFlow Lite?

A lightweight runtime for running machine learning models on mobile and embedded devices, optimized for inference with minimal latency.

We'll use the [TensorFlow Lite runtime](#) for Raspberry Pi, a simplified library for running machine learning models on mobile and embedded devices, without including all TensorFlow packages.

```
pip install tf-lite-runtime
```

```
tf-lite-runtime-2.14.0-cp311-cp311-manylinux_2_34_aarch64.whl
```

```
sudo reboot
```

Verify installation:

```
pip list | grep -E "(tf-lite-runtime)"
```

```
tf-lite-runtime 2.14.0.
```

Creating a **working directory** & get the **model**

```
Documents/  
├── TFLITE/  
│   └── IMG_CLASS/  
│       ├── models/  
│       │   ├── mobilenet_v2.tflite  
│       │   └── labels.txt  
│       └── images/  
│           └── test_image.jpeg
```

Go to the **models/** folder and get the Model and labels

```
wget
```

```
https://storage.googleapis.com/download.tensorflow.org/models/tflite\_11\_05\_08/mobilenet\_v2\_1.0\_224\_quant.tgz
```

```
tar xzf mobilenet_v2_1.0_224_quant.tgz
```

```
wget
```

```
https://raw.githubusercontent.com/Mjrovai/EdgeML-with-Raspberry-Pi/refs/heads/main/IMG\_CLASS/models/labels.txt
```

Verifying the Setup




```
import tf.lite_runtime.interpreter as tflite
import numpy as np
from PIL import Image

print("NumPy:", np.__version__)
print("Pillow:", Image.__version__)

# Try to create a TFLite Interpreter
model_path = "./models/mobilenet_v2_1.0_224_quant.tflite"
interpreter = tflite.Interpreter(model_path=model_path)
interpreter.allocate_tensors()
print("TFLite Interpreter created successfully!")
```

TFLite Interpreter



input_details

```
[{'name': 'input',  
  'index': 171,  
  'shape': array([ 1, 224, 224,  3], dtype=int32),  
  'shape_signature': array([ 1, 224, 224,  3], dtype=int32),  
  'dtype': numpy.uint8,  
  'quantization': (0.0078125, 128),  
  'quantization_parameters': {'scales': array([0.0078125], dtype=float32),  
  'zero_points': array([128], dtype=int32),  
  'quantized_dimension': 0},  
  'sparsity_parameters': {}}]
```


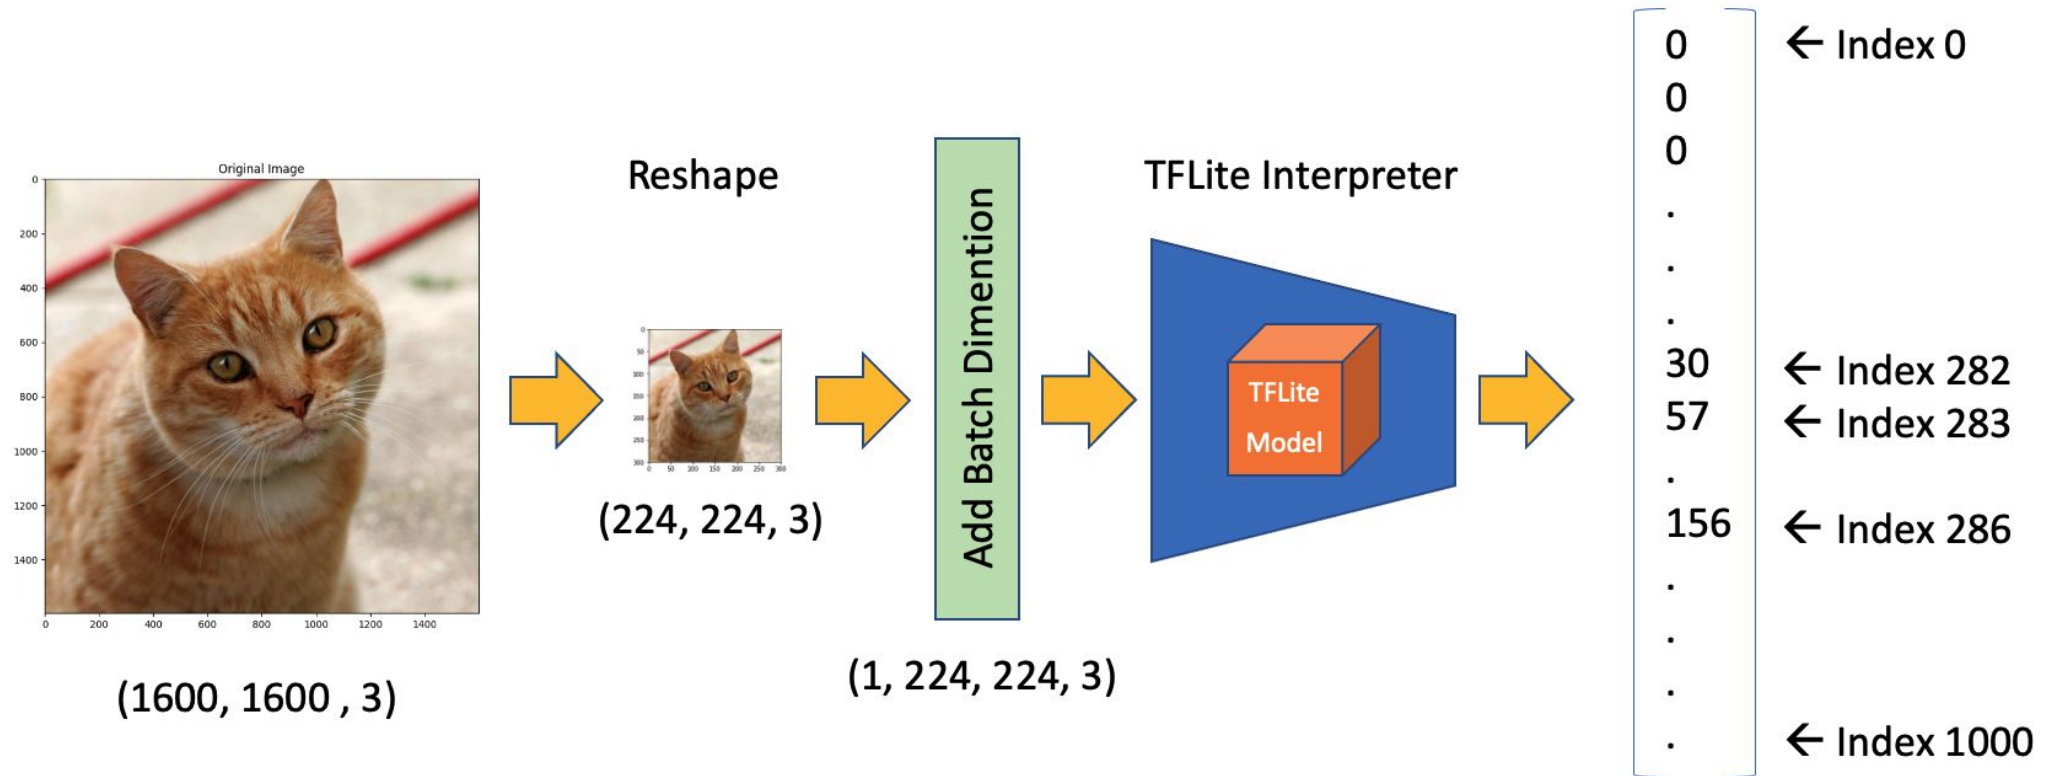
← Input Image Shape

output_details


```
[{'name': 'output',  
  'index': 172,  
  'shape': array([ 1, 1001], dtype=int32),  
  'shape_signature': array([ 1, 1001], dtype=int32),  
  'dtype': numpy.uint8,  
  'quantization': (0.09889253973960876, 58),  
  'quantization_parameters': {'scales': array([0.09889254], dtype=float32),  
  'zero_points': array([58], dtype=int32),  
  'quantized_dimension': 0},  
  'sparsity_parameters': {}}]
```

← Output model

Making **inferences** with MobileNet V2



Input Image



Pre-Process




Classification



Predict Label

10 Image Classification.ipynb



Making **inferences**: Static Images and Camera

[PREDICTION] [Prob]

| | |
|--------------|-------|
| tiger cat | : 37% |
| Egyptian cat | : 27% |
| tabby | : 16% |
| lynx | : 10% |
| carton | : 2% |



[PREDICTION] [Prob]

| | |
|-------------|-------|
| coffee mug | : 99% |
| cup | : 0% |
| whiskey jug | : 0% |
| teapot | : 0% |
| water jug | : 0% |



Questions?



Prof. Marcelo J. Rovai

rovai@unifei.edu.br



UNIFEI