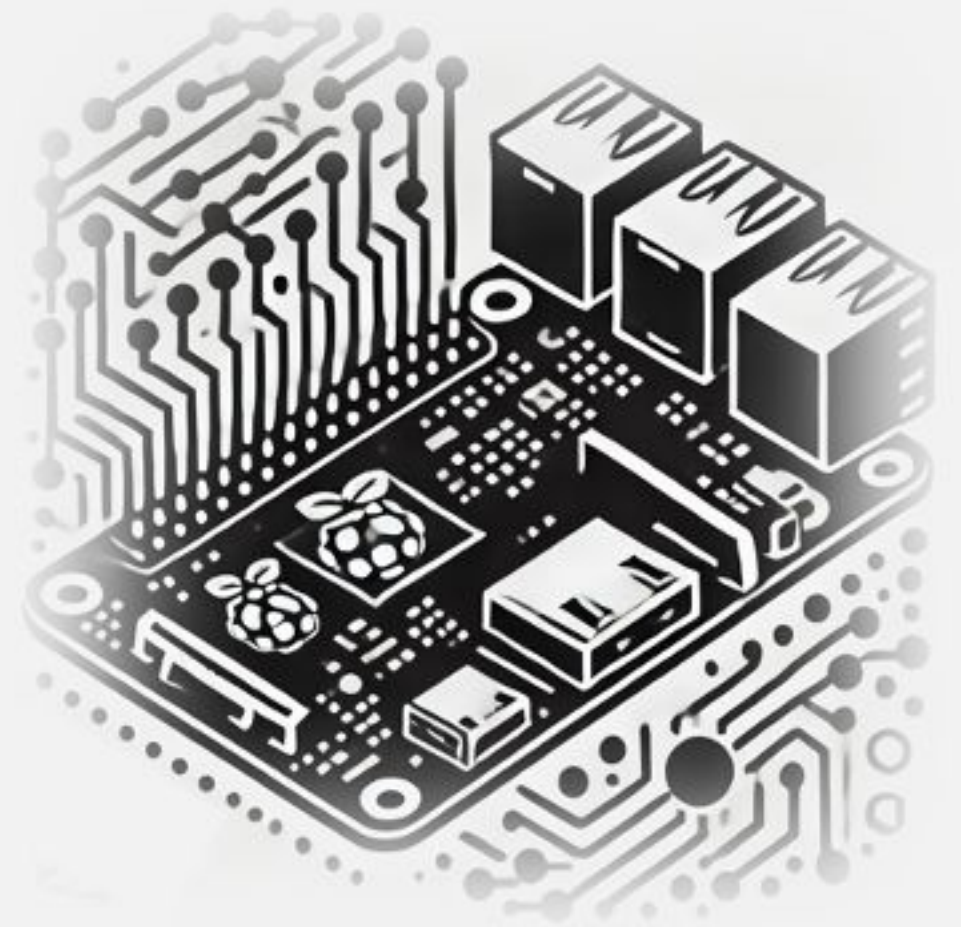


# IESTI05 – Edge AI

## Machine Learning System Engineering

---

### 5. Image Classification: Introduction to TFLite Runtime



# Computer Vision Recognition Tasks

## Image Classification (Multi-Class Classification)




Cat: 70%



Dog: 80%

## Object Detection Multi-Label Classification + Object Localization



# Computer Vision Recognition Tasks

## Instance Segmentation

Each **pixel** in an image IS CLASSIFIED into a predefined category.



## Pose Estimation

**Key points (or landmarks)** on the object, such as joints on a human body are detected



# Image Classification

# What is Image Classification?

## Definition

Image classification is a fundamental task in computer vision that involves **categorizing an image into one of several predefined classes**.

## Key Characteristics

- Assigns a **single label to the entire image**
- Based on visual content analysis
- Uses machine learning algorithms
- Mimics human visual perception



Cat: 70%



Dog: 80%



# Real-World Applications

## **Healthcare**

Medical image analysis, X-ray and MRI diagnostics

## **Agriculture**

Crop health monitoring, disease detection

## **Automotive**

Autonomous vehicles, road sign recognition

## **Security**

Surveillance systems, threat detection


## **Retail**

Visual search, inventory management


## **Environmental**

Satellite imagery analysis, monitoring

# Image Classification: Inference Pipeline




Input Image



Pre-Process

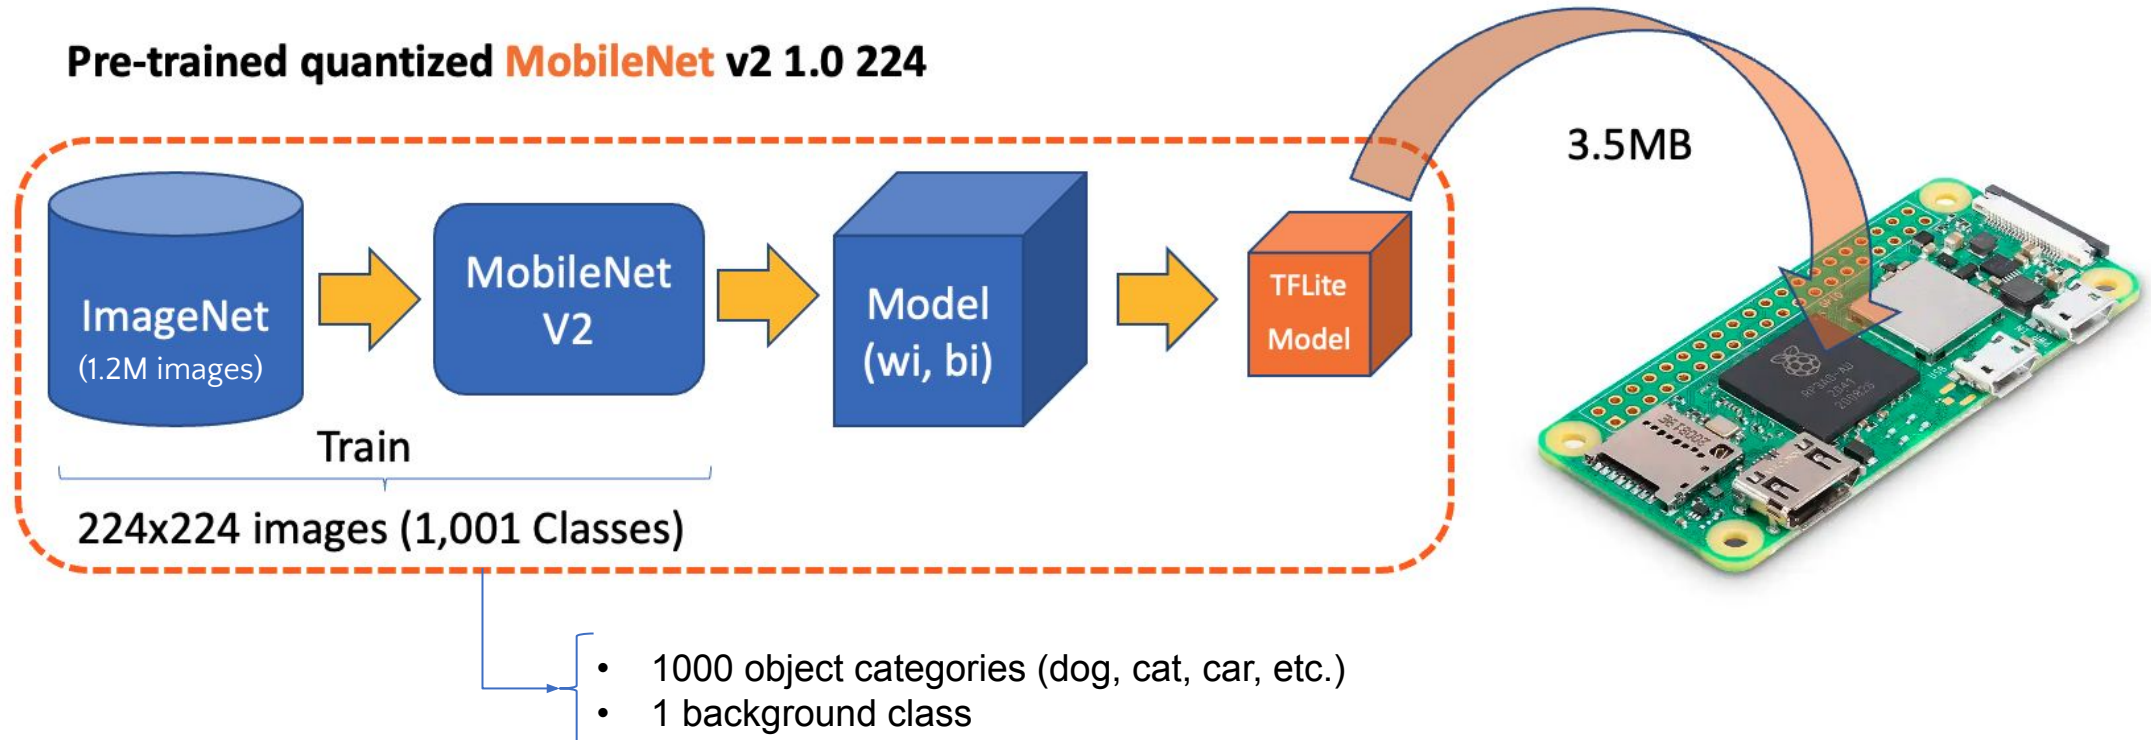


Classification  
Model




Predict Label

# The MobileNet V2 model





## TFLite Interpreter



### input\_details

```
[{'name': 'input',
  'index': 171,
  'shape': array([ 1, 224, 224,  3], dtype=int32),
  'shape_signature': array([ 1, 224, 224,  3], dtype=int32),
  'dtype': numpy.uint8,
  'quantization': (0.0078125, 128),
  'quantization_parameters': {'scales': array([0.0078125], dtype=float32),
  'zero_points': array([128], dtype=int32),
  'quantized_dimension': 0},
  'sparsity_parameters': {}}]
```

Input Image  
Shape Signature

Pixel data type

Quantization Parameters

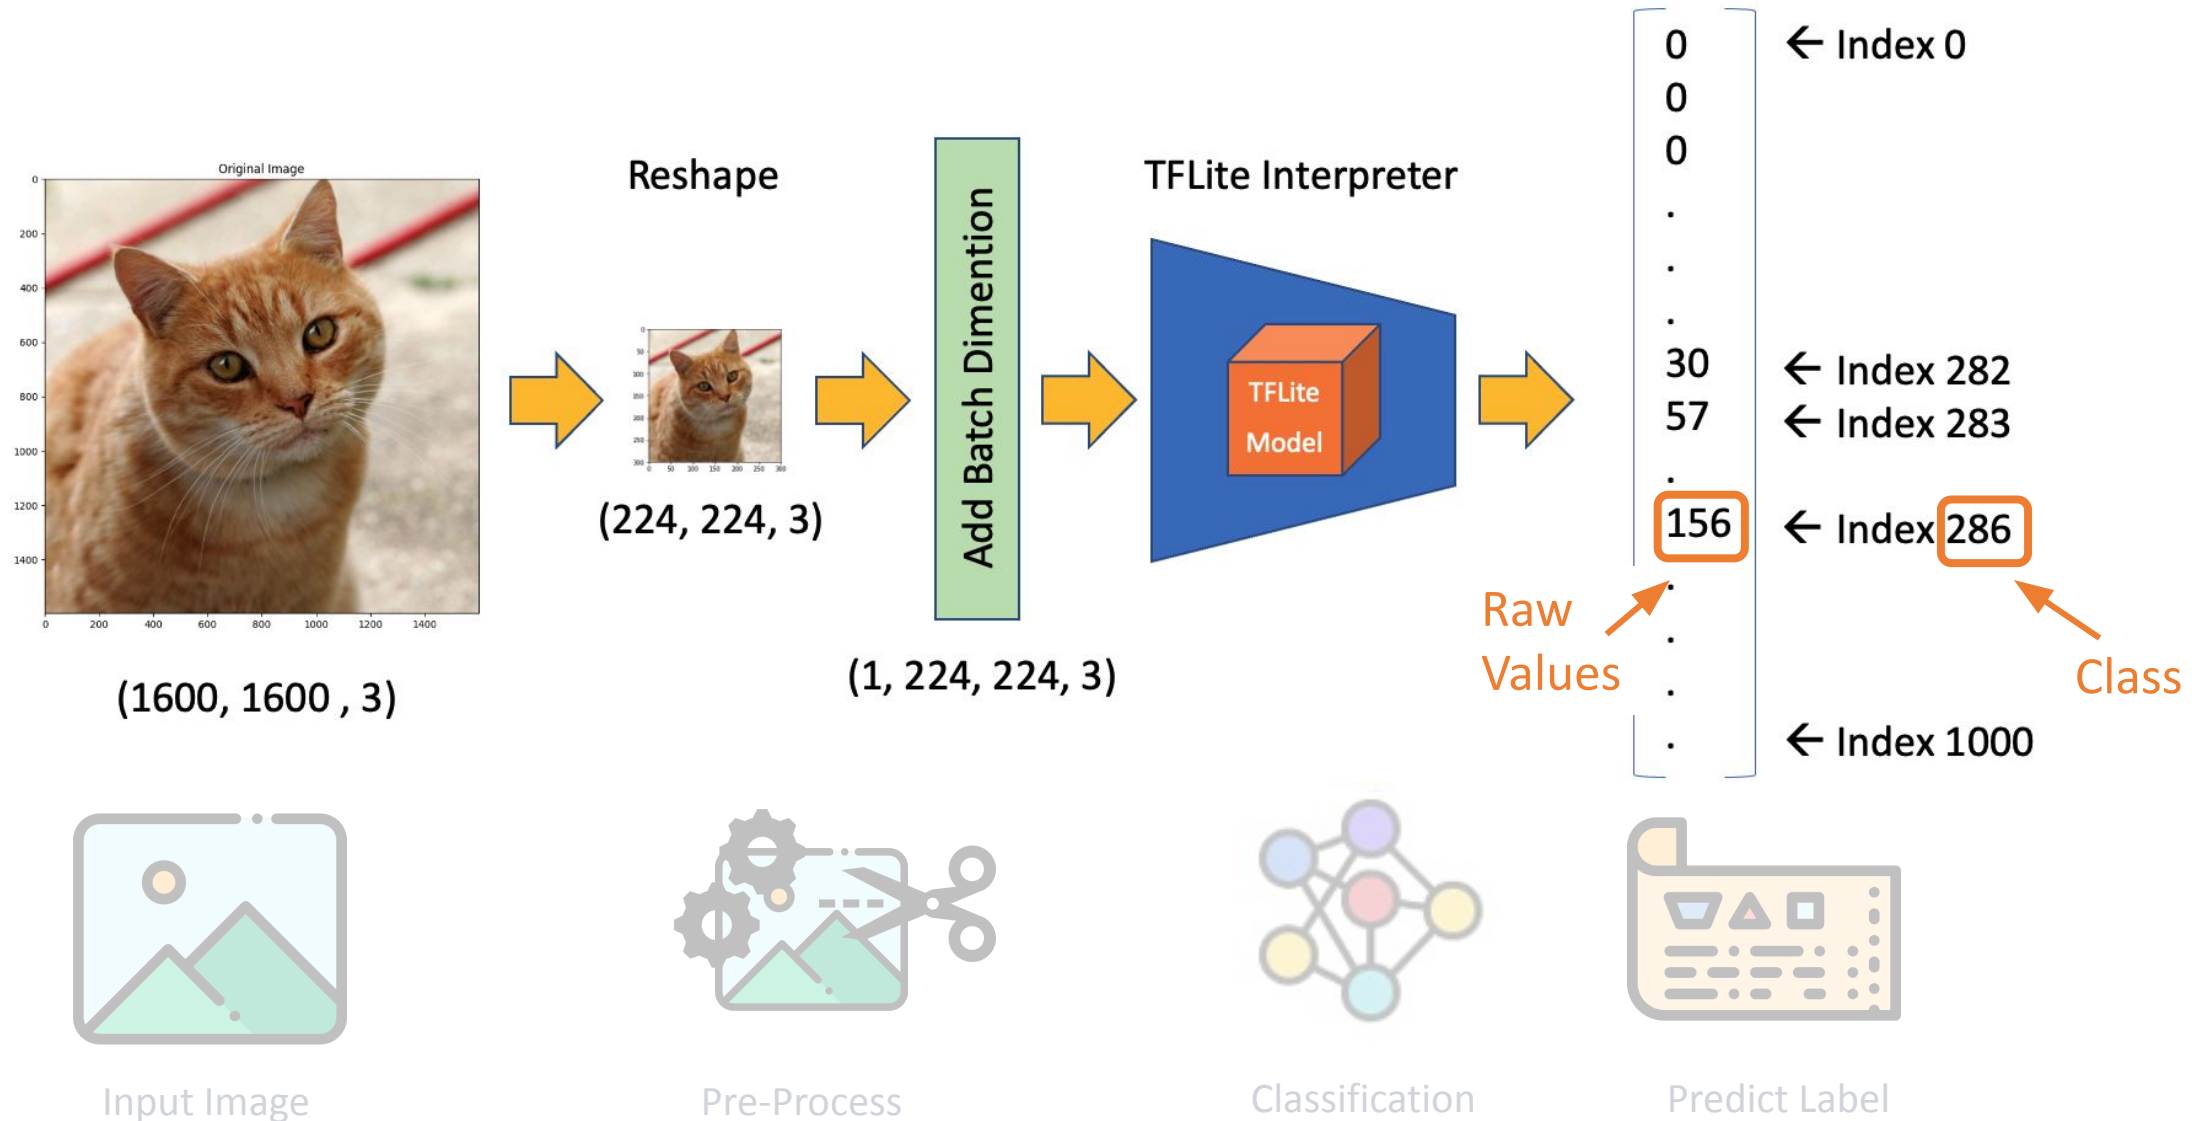
### output\_details

```
[{'name': 'output',
  'index': 172,
  'shape': array([ 1, 1001], dtype=int32),
  'shape_signature': array([ 1, 1001], dtype=int32),
  'dtype': numpy.uint8,
  'quantization': (0.09889253973960876, 58),
  'quantization_parameters': {'scales': array([0.09889254], dtype=float32),
  'zero_points': array([58], dtype=int32),
  'quantized_dimension': 0},
  'sparsity_parameters': {}}]
```

Output model

Quant Dim 0: One set of parameters for the entire output tensor

# Making **inferences** with MobileNet V2



# TFLite Runtime


# Setting up a Virtual Environment

**Activate** the environment:

```
source ~/tflite_env/bin/activate
```

To **exit** the virtual environment, use:

```
deactivate
```



# TensorFlow Lite Setup

## What is TensorFlow Lite?

A lightweight runtime for running machine learning models on mobile and embedded devices, optimized for inference with minimal latency.

We'll use the [TensorFlow Lite runtime](#) for Raspberry Pi, a simplified library for running machine learning models on mobile and embedded devices, without including all TensorFlow packages.

```
pip install tf-lite-runtime
```

```
tf-lite-runtime-2.14.0-cp311-cp311-manylinux_2_34_aarch64.whl
```

```
sudo reboot
```

Verify installation:

```
pip list | grep -E "(tf-lite-runtime)"
```

```
tf-lite-runtime 2.14.0.
```

# Creating a **working directory** & get the **model**

```
Documents/  
├── TFLITE/  
│   └── IMG_CLASS/  
│       ├── models/  
│       │   ├── mobilenet_v2.tflite  
│       │   └── labels.txt  
│       └── images/  
│           └── test_image.jpeg
```

Go to the **models/** folder and get the Model and labels

```
wget
```

```
https://storage.googleapis.com/download.tensorflow.org/models/tflite\_11\_05\_08/mobilenet\_v2\_1.0\_224\_quant.tgz
```

```
tar xzf mobilenet_v2_1.0_224_quant.tgz
```

```
wget
```

```
https://raw.githubusercontent.com/Mjrovai/EdgeML-with-Raspberry-Pi/refs/heads/main/IMG\_CLASS/models/labels.txt
```

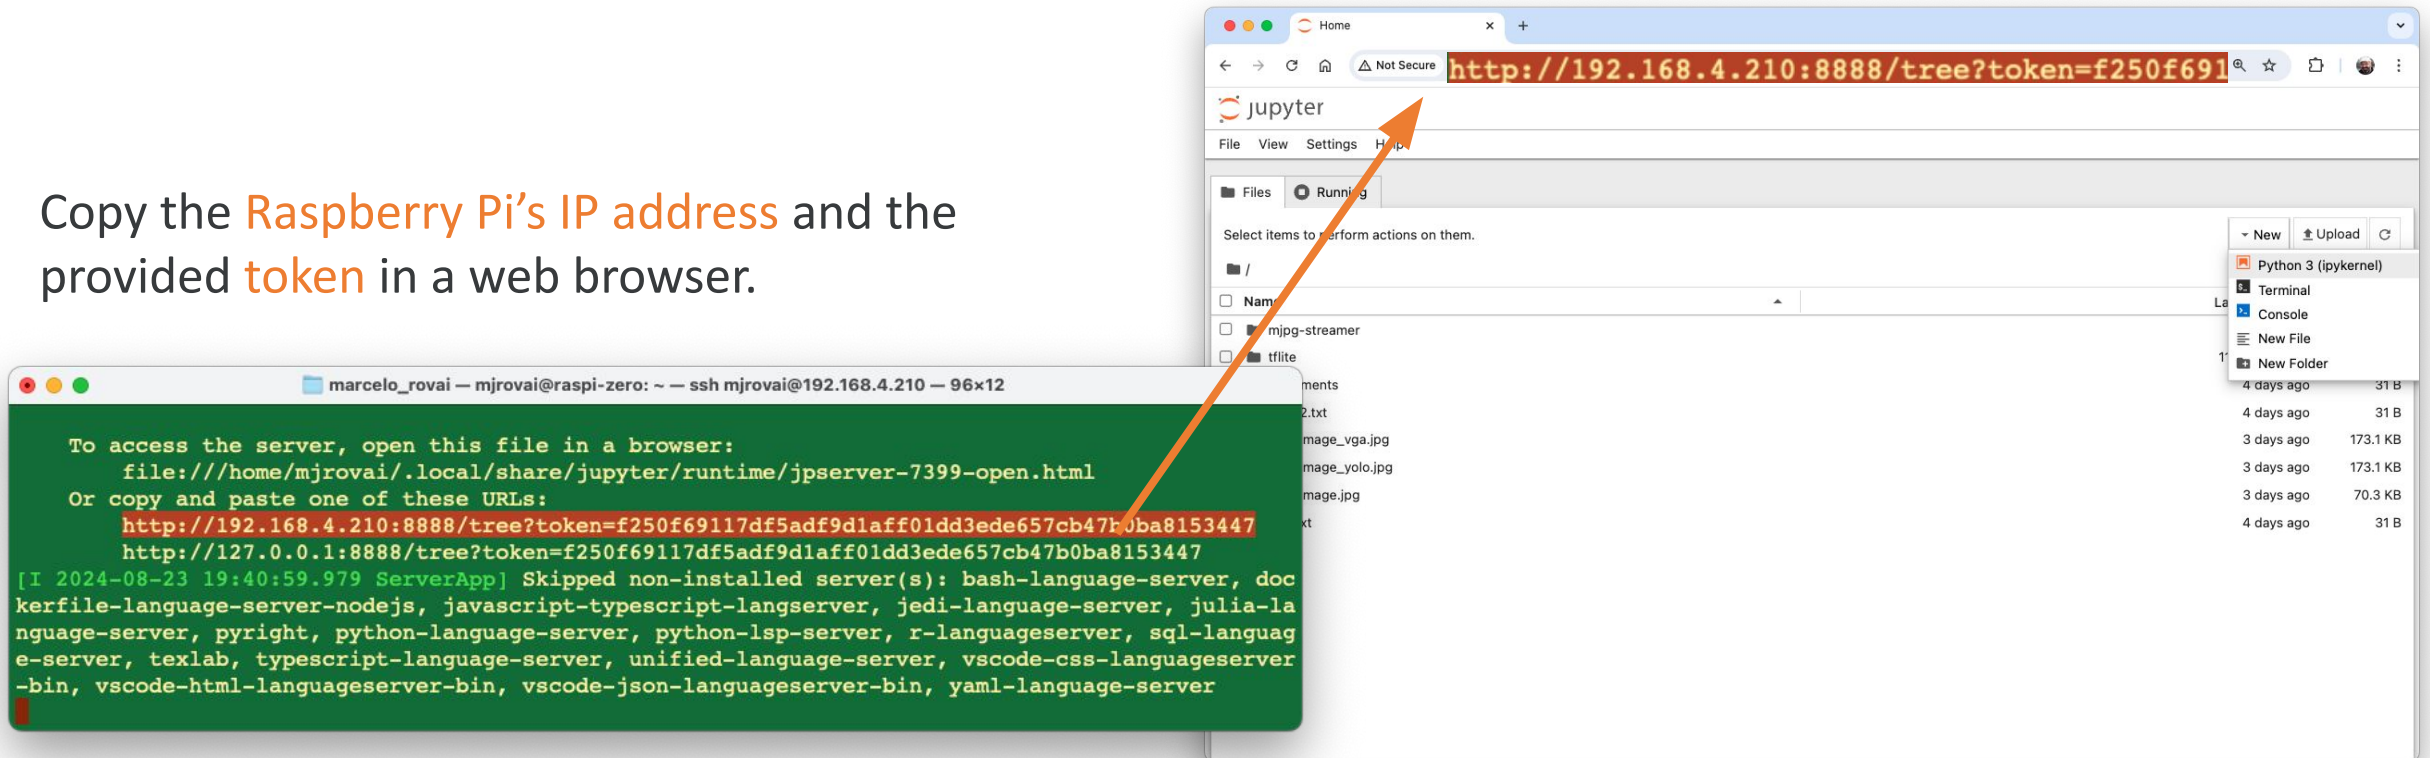


# Running up Jupyter Notebook



```
jupyter notebook --ip=[YOUR IP ADDRESS] --no-browser
```

Copy the **Raspberry Pi's IP address** and the provided **token** in a web browser.



# Verifying the Setup




```
import tf.lite_runtime.interpreter as tflite
import numpy as np
from PIL import Image

print("NumPy:", np.__version__)
print("Pillow:", Image.__version__)

# Try to create a TFLite Interpreter
model_path = "./models/mobilenet_v2_1.0_224_quant.tflite"
interpreter = tflite.Interpreter(model_path=model_path)
interpreter.allocate_tensors()
print("TFLite Interpreter created successfully!")
```

## Raspberry Pi Inference:

### 10\_Image\_Classification.ipynb



# Making inferences: Static Images and Camera

[PREDICTION]	[Prob]
--------------	--------

tiger cat	: 37%
Egyptian cat	: 27%
tabby	: 16%
lynx	: 10%
carton	: 2%



[PREDICTION]	[Prob]
--------------	--------

coffee mug	: 99%
cup	: 0%
whiskey jug	: 0%
teapot	: 0%
water jug	: 0%



# Optional


## Training:

[CNN model to classify the CIFAR-10 dataset](#)




## Raspberry Pi Inference:

[20\\_Cifar\\_10\\_Image\\_Classification.ipynb](#)





# Cifar-10 - Inference



[PREDICTION]	[Prob]
dog	: 45.0%
cat	: 28.7%
bird	: 9.3%
horse	: 5.5%
truck	: 4.2%

Inference time: 2.8ms



[PREDICTION]	[Prob]
cat	: 56.2%
dog	: 24.9%
frog	: 14.3%
deer	: 2.2%
bird	: 1.9%

Inference time: 2.8ms



[PREDICTION]	[Prob]
ship	: 88.7%
airplane	: 9.2%
deer	: 1.8%
bird	: 0.2%
automobile	: 0.0%

Inference time: 2.8ms



[PREDICTION]	[Prob]
automobile	: 99.7%
truck	: 0.3%
airplane	: 0.0%
ship	: 0.0%
bird	: 0.0%

Inference time: 2.8ms





# Questions?



Prof. Marcelo J. Rovai

[rovai@unifei.edu.br](mailto:rovai@unifei.edu.br)



**UNIFEI**