




IESTI05 – Edge AI



Machine Learning System Engineering

15. Text Generation – RNN The VerneBot





Unstructured Data - Text





Neural Network Architectures

Vibration
Analysis




Image
Classification




Text
Generation



Image
Generation




Large Language
Models- LLMs




MLP - Deep Neural Network




CNN - Convolutional NN




RNN - Recurrent NN (GRU/LSTM)




GAN - Generative Adversarial N.



AN - Attention (Transformers)





Neural Network Architectures

Vibration
Analysis




Image
Classification



Text
Generation



Image
Generation




Large Language
Models- LLMs



MLP - Deep Neural Network



CNN - Convolutional NN



RNN - Recurrent NN (GRU/LSTM)




GAN - Generative Adversarial N.



AN - Attention (Transformers)





20_JulesVerneBot_Generatin

https://colab.research.google.com/github/Mjrovai/TinyML4D/blob/main/WALC_2024/notebooks/20_JulesVerneBot_Generating_English_Texts_with...


CommandsCodeText

Connect T4

Generating Text with RNNs: The Jules Verne Bot

- By Marcelo Rovai @Nov24
- With code support from **Claude 3.5 Sonnet** and **ChatGPT 4o**







Project Gutenberg

DATASET

Total characters:

- 5,768,791

Unique characters:


- 123



1. 'A Journey to the Centre of the Earth'
2. 'In Search of the Castaways'
3. 'An Antarctic Mystery'
4. 'In the year 2889'
5. 'Around the World in Eighty Days'
6. 'Michael Strogoff'
7. 'Five Weeks in a Balloon'
8. 'The Mysterious Island'
9. 'From the Earth to the Moon'
10. 'Twenty Thousand Leagues under the Sea'

TOKENIZATION AND VOCABULARY

Conversion of text
into numeric tokens.



```
graph TD; A[Conversion of text into numeric tokens.] --> B[Character-level tokenization: 1 2 3 unique characters.]; B --> C[Example: "The Project" → [122 52 69 66 1 48 79 76 71 66 64 81].];
```

Character-level
tokenization: 1 2 3
unique characters.

Example: "The Project" →
[122 52 69 66 1 48 79 76 71 66 64 81].

TOKENIZATION

```
['\n', ' ', '!', '"', '#', '$', '%', '&', "'", '(', ')', '*', '+',  
, '-', '.', '/', '0', '1', '2', '3', '4', '5', '6', '7', '8',  
'9', ':', ';', '<', '=', '>', '?', 'A', 'B', 'C', 'D', 'E', 'F',  
'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S',  
'T', 'U', 'V', 'W', 'X', 'Y', 'Z', '[', ']', '_', 'a', 'b', 'c',  
'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p',  
'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', '{', '}', '~',  
'f', 'o', 'l', 'A', 'Æ', 'Ê', 'à', 'â', 'æ', 'ç', 'è', 'é', 'ê',  
'ë', 'ï', 'ñ', 'ô', 'û', 'ü', 'Œ', 'œ', 'ð', '—', '‘', '’', '“',  
'”', '•', 'ˆ', '™', '\uffeff'], dtype='<U1')
```

```
[31, 42, 40, 1, 30, 28, 46, 40, 48, 45, 45, 42, 0, 0, 36, 0, 0,  
31, 71, 1, 76, 65, 76, 77, 68, 71, 13, 0, 0, 48, 69, 57, 1, 70,  
71, 65, 76, 61, 1, 60, 61, 75, 76, 57, 75, 11, 1, 78, 65, 70, 60,  
71, 1, 60, 57, 1, 59, 65, 60, 57, 60, 61, 1, 72, 57, 74, 57, 1,  
71, 1, 32, 70, 63, 61, 70, 64, 71, 1, 41, 71, 78, 71, 11, 1, 61,  
70, 59, 71, 70, 76, 74, 61, 65, 1, 70, 71, 0, 76, 74, 61, 69, 1,  
60, 57, 1, 30, 61, 70, 76, 74, 57, 68, 1, 77, 69, 1, 74, 57, 72,  
57, 82, 1, 57, 73, 77, 65, 1, 60, 71, 1, 58, 57, 65, 74, 74, 71,  
11, 1, 73, 77, 61, 1, 61, 77, 1, 59, 71, 70, 64, 61])
```


TRAINING SEQUENCES



Goal: Predict the next character in a sequence.



Length of the sequence: 120 characters (~ paragraph).



Input 'Hello my nam'

Output 'ello my name'.

EMBEDDING


Each character is represented
as a vector of 256
dimensions.



Embedding captures
relationships between
characters in dense vectors.

EMBEDDING

3D Projection of Character Embeddings Using UMAP



MODEL ARCHITECTURE

Embedding Layer:
Converts characters into
dense vectors.




RNN/GRU Layer (1024
units): Learn from
sequences.




Dense layer: Generates
probabilities for each
character (123).



Deep Learning models (or artificial neural networks)

Recurrent Neural Networks (RNNs): Designed for **sequential data like time series or text**, these networks use their internal state (memory) to process sequences of inputs.





Gated Recurrent Unit (GRU)



Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(1, 120, 256)	31,488
gru_3 (GRU)	(1, 120, 1024)	3,938,304
dense_3 (Dense)	(1, 120, 123)	126,075

Total params: 4,095,867 (15.62 MB)
 Trainable params: 4,095,867 (15.62 MB)
 Non-trainable params: 0 (0.00 B)

RNN MODEL (RECURRENT)

MODEL TRAINING

Loss Function: Categorical
Sparse Crossentropy


Optimizer: Adam

Epochs: 30

batch size: 128

buffer size: 10,000 (for
dataset shuffling)


Monitoring loss reduction
over time.



(33 minutes for training)

Google
colab

Tesla T4 GPU



TEXT GENERATION

The template generates
text character by character
from an initial text:

"THE FLYING
SUBMARINE".

Temperature controls
randomness (0.5 for
predictable, 1.0 for
creative text).

Generated text with temperature 0.5:

THE FLYING SUBMARINE

CHAPTER 100 VENTANTILE

This eBook is for the use of anyone anywhere in the United States and most other parts of the earth and miserable eruptions. The solar rays should be entirely under the shock of the intensity of the sea. We were all sorts. Are we to prepare for our feelings?"

"I can never see them a good geographer," said Mary.

"Well, then, John, for I get to the Pampas, that we ought to obey the same time. In the country of this latitude changed my brother, and the Nautilus floated in a sea which contained the rudder and lower colour visibly. The loiter was a fatalint region the two scientific discoverers. Several times turning toward the river, the cry of doors and over an inclined plains of the Angara, with a threatening water and disappeared in the midst of the solar rays.

The weather was spread and strewn with closed bottoms which soon appeared that the unexpected sheets of wind was soon and linen, and the whole seas were again landed on the subject of the natives, and the prisoners were successively assuming the sides of this agreement for fifteen days with a threatening voice.

The clouds had disappeared to the ground, and the river and the ship's conditions of the ship's course was still standing, but in his daring explorers, and the sailor thought for the sudden discovery.

"There are no trees, and a half an hour or the sun will soon be carried off they were bringing a special track."

"As you see, my dear Helena, who was the matter of despair?" cried the captain.

"The raft we had done now, captain, and I have a reasonable face of the shipwrecked creek, had been discovered at the entrance of the world."

CHALLENGES AND LIMITATIONS

*Limited **context window** (120 characters).*

Difficulty in maintaining coherence in long texts.

Character-level modeling vs. word-level modeling.

CONNECTING WITH MODERN LANGUAGE MODELS



Our Model (VerneBot) :

- Training data: **5.8 million characters (bytes)** (10 books).
- **4 million parameters**,
- Character-level tokenization (**123**)
- **RNN** architecture.




Open AI GPT-3 (2020):

- Training data: **45 Trillion bytes** (text)
- **175 billion parameters**,
- Subword tokenization (**2,048 tokens**),
- **Transformer** Architecture.



Modern models handle long-range dependencies better.



- **Architecture:** Transformer – 3.8 Billion Parameters


Inputs: Text.

Context length: 128k tokens

GPU: 512 H100-80G

Training time: 7 days

Training data: 3.3 Trillion tokens**



~ 350 pages

~ 300 words/page

1 word = ~ 1.4 token

** Equivalent to 23 million books, that is:
17% of all the books in the world

Aspect	Jules Verne Bot	GPT-3 (2020)	Phi-3-mini (2024)
Architecture	RNN (GRU)	Transformer	Transformer
Parameters	4 million	175 billion	3.8 billion
Training Data	5.8M characters	45TB text	3.3T tokens
Context Length	120 characters	2,048 tokens	128,000 tokens
Tokenization	Character-level	Subword (BPE)	Subword
Training Time	33 minutes	Months	7 days
GPU Requirements	1 Tesla T4	Thousands of GPUs	512 H100 GPUs

Scale Comparison

Questions?

Prof. Marcelo J. Rovai

rovai@unifei.edu.br



UNIFEI